



■ Real-time event streaming with Python

Dave Klein - Senior Developer Advocate



Stream Processing






Apache Kafka



Apache Pulsar



Red Panda




Amazon Kinesis



... and more



Apache Flink



ksqlDB



Kafka Streams




... and more



Apache Kafka

A Primer

■ A Log (Topic)




Immutable
Append-only
Ordered
Durable

```
"timestamp": "2023-04-27T18:30:01.223Z",  
"key": "driver-0032",  
"value": {  
  "eventType": "delivery-update",  
  "lat": 51.12321322,  
  "lon": 14.132131212,  
  "speed": 35,  
  "deliveryId": "pizza-my-heart-123213"  
}
```




The diagram illustrates a Producer-Consumer queue system. At the top, a green rounded rectangle labeled 'Producer' has a curved arrow pointing to the 9th slot of a horizontal queue. The queue is a row of 11 boxes, numbered 1 to 11 below them. Slots 1 through 9 contain a document icon, while slots 10 and 11 are empty. A second curved arrow points from the 4th slot of the queue to a pink rounded rectangle at the bottom labeled 'Consumer'. Below the 'Consumer' label, the text 'Committed Offset: 3' is displayed.

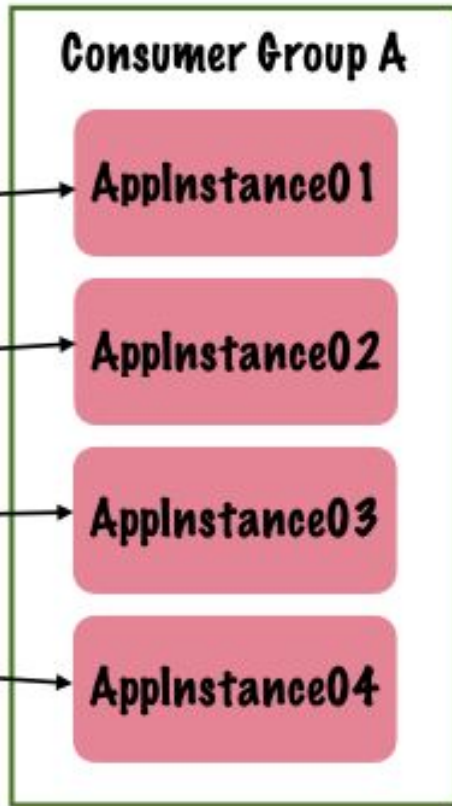

Producer



Consumer

Committed Offset: 3





- Python Kafka Clients

- kafka-python

<https://github.com/dpkp/kafka-python>

- confluent-kafka

<https://github.com/confluent-inc/confluent-kafka-python>

- aiokafka

<https://github.com/aio-labs/aiokafka>

■ Producing data to a Kafka topic

```
from confluent-kafka import Producer

producer = Producer({"bootstrap.servers": "localhost:29092"})

# some external system is feeding us delivery updates
del_update = recieve_delivery_update()

producer.produce("food-delivery", del_update, del_update["driver_id"])

producer.flush()
```

● Consuming data from a Kafka topic

```
from confluent-kafka import Consumer
import json

consumer = Consumer({"bootstrap.servers": "localhost:29092",
                    "group.id": "deliveries"})
consumer.subscribe(["food-delivery"])

while True:
    event = consumer.poll(1.0)
    if event is not None:
        val = event.value().decode("utf-8")
        del_update = json.loads(val)
        # do something with the event data
```



Stream Processing

■ Stream Processing – Stateless Operations

Filter

```
producer = Producer(...)
consumer = Consumer(...)

consumer.subscribe(["input_topic"])

while True:
    event = consumer.poll(1.0)
    if check_predicate(event):
        producer.produce("output_topic", event.value(), event.key())
```

■ Stream Processing – Stateless Operations

Map

```
producer = Producer(...)
consumer = Consumer(...)
consumer.subscribe(["input_topic"])

def transform_event(event):
    # some process we want done on each event

while True:
    event = consumer.poll(1.0)
    if event is not None:
        transformed_value = transform_event(event)
        producer.produce("output_topic", transformed_value, event.key())
```

● Stream Processing – Stateful Operations

Count

. . .

```
counts = {}  
while True:  
    event = consumer.poll(1.0)  
    if event.key() in counts:  
        counts[event.key()] += 1  
    else:  
        counts[event.key()] = 1  
  
    producer.produce("output", counts[event.key()], event.key())
```


■ Stream Processing – Stateful Operations

Sum

```
. . .
```

```
sums = {}
```

```
while True:
```

```
    event = consumer.poll(1.0)
```

```
    if event.key() in sums:
```

```
        sums[event.key()] += event.value()
```

```
    else:
```

```
        sums[event.key()] = event.value()
```

```
    producer.produce("output", sums[event.key()], event.key())
```

■ Stream Processing – Stateful Operations

Join

```
# a_dict, b_dict, c_dict, and consumers declared above
while True:
    a_event = consumer_a.poll(1.0)
    a_dict[a_event.key()] = a_event
    if a_event.key() in b_dict and a_event.key() not in c_dict:
        c_event = join_func(a_event, b_dict[a_event.key()])
        c_dict[a_event.key()] = c_event
    b_event = consumer_b.poll(1.0)
    b_dict[b_event.key()] = b_event
    if b_event.key() in a_dict and b_event.key() not in c_dict:
        c_event = join_func(a_dict[b_event.key()], b_event)
        c_dict[b_event.key()] = c_event
    producer.produce("output", key=c_event.key(), value=c_dict[c_event.key()])
```



Demo

[**https://github.com/daveklein/top-tweeters**](https://github.com/daveklein/top-tweeters)

Top Tweeters for PyCon Italia 2023

User	Count
Saurav Jain (Open Source + Communities)	20
Marlene Mhangami	14
Ester 🇺🇦	14
🇮🇳🇧🇩🇮 fundor333@mastodon.social 🌐🏠	10
Cheuk Ting Ho at #PyCon IT	10
Danica Fine	9
Matteo Benci	9
Alessia Marcolini	7
Marcelo Trylesinski	7
Fiorella De Luca	7
Paolo Castagna	6
Valerio Maggio @leriomaggio@mastodon.social	6

■ Stream Processing Challenges

State

Scale

■ Python Streaming Clients

■ quixstreams

<https://github.com/quixio/quix-streams>

■ faust

<https://github.com/faust-streaming/faust>

● do-it-yourself

Processing data with Quix Streams

```
import quixstreams as qx

client = qx.KafkaStreamingClient("localhost:29092")
consumer = client.get_topic_consumer("food-delivery", consumer_group="eta_calc")
producer = client.get_topic_producer("food-delivery-with-eta")

def received_handler(stream: qx.StreamConsumer, df: pd.DataFrame):
    df["ETA"] = calc_eta(df[["lat", "lon"]], stream.stream_id)
    producer.timeseries.publish(df)

def on_stream_received_handler(stream_consumer: qx.StreamConsumer):
    stream_consumer.timeseries.on_data_received = received_handler

consumer.on_stream_received = on_stream_received_handler

qx.App.run()
```

Thank you



dave@quix.io | forum.quix.io